

Nishant Sharma

+1 (646) 772-9142 | hellonishantsh@gmail.com | [linkedin/in/nishantsh20](https://www.linkedin.com/in/nishantsh20) | hellonish.dev
Initial OPT (STEM eligible) | New York, NY

Education

*Master of Science in Computer Engineering, **New York University***

Graduated: May 2026

Served as ML course assistant. Worked on language modelling for reasoning tasks and fine-tuning LLMs.

Work Experience

Founding AI Engineer, [Felix Labs](#), New York, NY

Sep 2025 to Present

- Produced open-source singularity agent CLI and Web developed in Python/FastAPI and NextJS.
- Developed crash-resilient infrastructure for long-horizon tasks using langgraph and Redis/ARQ for queuing and Supabase for persistence, live over AWS.
- Engineered a multi-agent system for research and search workflow with sandbox execution for zero-trust tasks.

Software Engineer, [Ingelt](#), New Delhi, DL

Sep 2023 to Aug 2024

- Built FlaskAPI and Node.js services for a multi-tenant EdTech SaaS serving 1,000+ students across 50+ organizations.
- Developed WebSocket chat feature and LLM-assist workflows using Redis pub/sub, per-org budgets, and rate-limits integrating to NextJS app.
- Cut API latency through S3 presigned uploads and celery-based async processes.

Software Engineer Intern (Backend), [Macverin](#), New Delhi, DL

Jul 2022 to Aug 2023

- Built REST APIs, and business layer operations for 10+ client products using Node.js, Django, including third-party integrations listed - Auth0, Google APIs, Stripe, and Twilio.
- Introduced Docker-based deployments and Jenkins CI/CD for high-value client systems across AWS, GCP, and Hostinger, improving deployment reliability.

Projects

AI Job Search Management Web Platform - [ineedajob.pro](#)

Skills: Python, FastAPI, NextJS, PostgreSQL, Redis, AWS

- Developed a job applications tracking and reachout platform in NextJS and FastAPI live over AWS with LangGraph AI workflows and job offloading to workers via a queue system implemented using Redis and built chrome extension for UX.

Multimodal Image Captioning MLOps - [blog](#)

Skills: LLaVA, QLoRA, PyTorch, FastAPI, Kubernetes, MLflow, MinIO, Argo Workflows, OpenStack

- Fine-tuned LLaVA-1.5-7B with 4-bit QLoRA on validated InstaCities1M image-caption pairs and built an E2E MLOps platform with MLflow artifact management, Kubernetes deployment environment, and Prometheus-based inference monitoring on Chameleon.

SLM Fine-tuning for Math Reasoning - [blog](#)

Skills: LLMs, Fine-tuning, Pytorch, Tensorflow, LoRA, SFT, Exploratory Data Analysis

- Fine-tuned Phi-2 (2.7B) on GSM8K with 4-bit QLoRA, Built RoBERTa-large outcome-verifier and pairwise-ranker models; concluded ranker-weighted voting nearly matched majority vote (68.16% vs. 68.61%), identifying generator stability as primary bottleneck.